

Intermediate Level C++

C++11: Move Semantics and rvalue References

Yong Zhang

Objectives

In this chapter you will learn:

- Move Semantics Introduction
- lvalue and rvalue revision
- rvalue References
- Move semantics implementation
- std::move
- Perfect forwarding

Move Semantics

- A new feature introduced in C++ which is an addition to copying via copy constructor or copy assignment operator.
- Classes can have move constructor and move assignment operators.

```
class JetPlane
{
public:
    JetPlane();

    JetPlane(const JetPlane&);
    JetPlane& operator=(const JetPlane&);

    JetPlane(JetPlane&&);
    JetPlane& operator=(JetPlane&&);

};
```

Move Semantics

- **Benefits of using move semantics:**
 - Better performance
 - Better code
 - Better support for exclusive resource ownership
- **How to implement move semantics:**
 - Using rvalue references

Lvalue and Rvalue

- Every expression in C++ is either an lvalue or an rvalue.
- What are **lvalues** and **rvalues**:

lvalue

- ✓ Has a name
- ✓ Can have address taken

var

*ptr

arr[n]

++x

rvalue

- ✓ Doesn't have a name
- ✓ Can't have address taken
- ✓ this pointer

a * b

x++

string("abc")

Lvalue and Rvalue (continued)

- The old definitions of lvalue and rvalue are not accurate:
 - Lvalues: An expression that may appear on the **left** or on the **right** hand side of an assignment.
 - Rvalues: An expression that can only appear on the **right** hand side of an assignment.

```
void func(const int * pi, const int & ri) {  
  
    *pi=7;//compilation error, *pi is const  
  
    ri=8; //compilation error, ri is const  
  
}
```

7==x;

Lvalue and Rvalue (continued)

- Lvalue or rvalue based on **return type** for function or operator calls:

```
vector<int> v;
v[0];           // lvalue
v.size();        // rvalue

string s;
s + "abc";      // rvalue
```

Lvalue and Rvalue (continued)

- Both lvalues and rvalues can have a const attribute, which results in four types of expressions:

```
string f() { return string("F"); }
const string g() { return string("G"); }
JetPlane jet;
const int max_power_level = 100;

jet;                      // lvalue
max_power_level;          // const lvalue
f();                      // rvalue
g();                      // const rvalue

cout << jet.model().append("_RR").size() << endl;
```

Lvalue and Rvalue (continued)

- Initializing references using lvalues and rvalues:

```
JetPlane jet;
JetPlane& jet_ref = jet;

const JetPlane grounded_jet;
JetPlane& jet_ref2 = grounded_jet;      // doesn't compile

JetPlane& jet_ref3 = JetPlane();        // doesn't compile

auto make_const_jet = []() -> const JetPlane { return JetPlane(); };
JetPlane& jet_ref4 = make_const_jet(); // doesn't compile
```

```
const JetPlane grounded_jet;
const JetPlane& jet_ref2 = grounded_jet;      // OK

const JetPlane& jet_ref3 = JetPlane();        // OK

auto make_const_jet = []() -> const JetPlane { return JetPlane(); };
const JetPlane& jet_ref4 = make_const_jet(); // OK
```

```
const JetPlane& jet_ref4 = make_const_jet(); // OK

jet_ref4; // the expression jet_ref4 is an lvalue - it has a name!
```

Lvalue and Rvalue (continued)

- **Alternate C++ definition for lvalues and rvalues:**

- Lvalues: An expression that refers to a memory location and allows us to take the address of that memory location via the **&** operator.
- Rvalues: An expression that is not an lvalue, or refers to a **temporary object**.

```
// lvalues:  
//  
int i = 42;  
i = 43; // ok, i is an lvalue  
int* p = &i; // ok, i is an lvalue  
int& foo();  
foo() = 42; // ok, foo() is an lvalue  
int* p1 = &foo(); // ok, foo() is an lvalue  
  
// rvalues:  
//  
int foobar();  
int j = 0;  
j = foobar(); // ok, foobar() is an rvalue  
int* p2 = &foobar(); // error, cannot take  
j = 42; // ok, 42 is an rvalue
```

Rvalue References

- **Rvalue References:** if X is any type, then **X&&** is called an rvalue reference to X.
- Lvalue and rvalue references are distinct types, but they behave similarly.
- The difference in initialization is that an rvalue reference can only be initialized with a **non-const rvalue**:

```
JetPlane&& jet_ref9 = JetPlane();           // OK

JetPlane jet;
JetPlane&& jet_ref10 = jet;                  // doesn't compile

const JetPlane grounded_jet;
JetPlane&& jet_ref11 = grounded_jet;        // doesn't compile

JetPlane&& jet_ref12 = make_const_jet(); // doesn't compile
```

Rvalue References (continued)

- In C++11, there are **four** reference-based overloads instead of 2.

```
void f(JetPlane& plane);
void f(const JetPlane& plane);
void f(JetPlane&& plane);
void f(const JetPlane&& plane);
```

- When it comes to function overload resolution, lvalues prefer old-style lvalue references, whereas rvalues prefer the new rvalue references:

```
void foo(X& x); // lvalue reference overload
void foo(X&& x); // rvalue reference overload

X x;
X foobar();

foo(x); // argument is lvalue: calls foo(X&)
foo(foobar()); // argument is rvalue: calls foo(X&&)
```

Rvalue References (continued)

- More on overload resolution:

```
JetPlane jet;
f(jet);           // f(JetPlane&)

const JetPlane grounded_jet;
f(grounded_jet);    // f(const JetPlane&)

f(JetPlane());      // f(JetPlane&&)

auto make_const_jet = []() -> const JetPlane { return JetPlane(); };
f(make_const_jet()); // f(const JetPlane&&)
```

f(make_const_jet())  f(const JetPlane&)

Move Semantics Implementation

- Rvalue references allow a function to branch at **compile time** via overload resolution.
- This kind of overload should occur only for **copy constructors** and **assignment operators**, for the purpose of achieving move semantics.

```
X& X::operator=(X const & rhs); // classical implementation
X& X::operator=(X&& rhs)
{
    // Move semantics: exchange content between this and rhs
    return *this;
}
```

Move Semantics Implementation (continued)

```
A(const A& rhs);
A(A&& rhs);

struct A
{
    A()
    {
        cout << "A's constructor" << endl;
    }
    A(const A& rhs)
    {
        cout << "A's copy constructor" << endl;
    }
};

vector<A> v;
cout << "==> push_back A():" << endl;
v.push_back(A());
cout << "==> push_back A():" << endl;
v.push_back(A());
```

```
==> push_back A():
A's constructor
A's copy constructor
==> push_back A():
A's constructor
A's copy constructor
A's copy constructor
```

Move Semantics Implementation (continued)

```
A(A&& rhs)
{
    cout << "A's move constructor" << endl;
}

vector<A> v;
cout << "==> push_back A():" << endl;
v.push_back(A());
cout << "==> push_back A():" << endl;
v.push_back(A());
```

```
==> push_back A():
A's constructor
A's move constructor
==> push_back A():
A's constructor
A's move constructor
A's move constructor
```

Move Semantics Implementation (continued)

- Another example of why move semantics is better:

```
X& X::operator=(X const & rhs)
{
    // [...]
    // Destruct the resource that m_pResource refers to.
    // Make a clone of what rhs.m_pResource refers to, and
    // attach it to m_pResource.
    // [...]
}

X foo();
X x;
// perhaps use x in various ways
x = foo();

X& X::operator=(<mystery type> rhs)
{
    // [...]
    // swap this->m_pResource and rhs.m_pResource
    // [...]
}
```

Move Semantics Implementation (continued)

```
class A
{
    double _d;
    int* _p;
    string _str;
public:
    A(A&& rhs) : _d(rhs._d), _p(rhs._p), _str(move(rhs._str))
    {
        rhs._p = nullptr;
        rhs._str.clear();
    }
    A& operator=(A&& rhs)
    {
        delete _p;

        _d = rhs._d;
        _p = rhs._p;
        _str = move(rhs._str); // careful!
        rhs._p = nullptr;
        rhs._str.clear();

        return *this;
    }
};
```



- Tell the compiler to choose a move overload when you
 - have a named rvalue reference, or
 - an lvalue that you know isn't going to be used anymore
- it does not move anything; it just turns an lvalue into an rvalue, so that the move constructor is invoked.
- Without move, the compiler would choose the copy overload.

```
A a;  
v.push_back(move(a)); // move overload used
```

Move Semantics Implementation – An Example

```
1  class ArrayWrapper
2  {
3      public:
4          ArrayWrapper (int n)
5              : _p_vals( new int[ n ] )
6              , _size( n )
7      {}
8      // copy constructor
9      ArrayWrapper (const ArrayWrapper& other)
10         : _p_vals( new int[ other._size ] )
11         , _size( other._size )
12     {
13         for ( int i = 0; i < _size; ++i )
14         {
15             _p_vals[ i ] = other._p_vals[ i ];
16         }
17     }
18     ~ArrayWrapper ()
19     {
20         delete [] _p_vals;
21     }
22     private:
23     int *_p_vals;
24     int _size;
25 }
```

Move Semantics Implementation – An Example (continued)

```
// move constructor
ArrayWrapper (ArrayWrapper&& other)
    : _p_vals( other._p_vals )
    , _size( other._size )
{
    other._p_vals = NULL;
    other._size = 0;
}
```

Move Semantics Implementation – An Example (continued)

```
1 #include <utility> // for std::move
2
3 // move constructor
4 ArrayWrapper (ArrayWrapper&& other)
5     : _p_vals( other._p_vals )
6     , _metadata( std::move( other._metadata ) )
7 {
8     other._p_vals = NULL;
9 }
```

- Demo: RvalueReferences

Perfect Forwarding

- The problem:

```
auto one_func(T arg1, T arg2)
{
    another_func(arg1, arg2);
}
```

- lengthy code:

```
unique_ptr<vector<Point>> p_points(new vector<Point>(10));
```

- Forward argument by value:

```
template<typename T, typename Arg>
unique_ptr<T> make_unique(Arg arg)
{
    return unique_ptr<T>(new T(arg));
}
```

- Forward argument by reference:

```
template<typename T, typename Arg>
unique_ptr<T> make_unique(Arg& arg)
{
    return unique_ptr<T>(new T(arg));
}
```

```
make_unique<vector<int>>(10); // can't convert argument from int to int&
```

Perfect Forwarding

- Is this the solution?

```
template<typename T, typename Arg>
unique_ptr<T> make_unique(const Arg& arg)
{
    return unique_ptr<T>(new T(arg));
}

template<typename T, typename Arg>
unique_ptr<T> make_unique(Arg& arg)
{
    return unique_ptr<T>(new T(arg));
}

int a = 10;
make_unique<vector<int>>(a);    // OK, Arg& overload
make_unique<vector<int>>(10);   // OK, const Arg& overload
```

Perfect Forwarding

- The solution:

```
template <typename T, typename T1, typename T2>
unique_ptr<T> make_unique(T1&& arg1, T2&& arg2)
{
    return unique_ptr<T>(new T(forward<T1>(arg1), forward<T2>(arg2)));
}
```

```
template<typename T, typename... Args>
unique_ptr<T> make_unique(Args&&... args)
{
    return unique_ptr<T>(new T(forward<Args>(args)...));
}
```



```
#include <utility>
```